

Monitoring for Developers with Prometheus and Grafana



Isaac Rubio Torres / July 2018


Monitoring for Developers with Prometheus and Grafana

- 1 Prometheus Manifesto
- 2 Setup
- 3 How to...
- 4 Prometheus works for Developers (and Ops)

Monitoring for Developers with Prometheus and Grafana

- 1 **Prometheus Manifesto**
- 2 Setup
- 3 How to...
- 4 Prometheus works for Developers (and Ops)

Monitoring



Prometheus is a Monitoring System and Time Series Database



Prometheus is an opinionated solution
for

instrumentation, collection, storage
querying, alerting, dashboards, trending

Prometheus values ...

operational systems monitoring (not only) for the cloud

over

raw logs and events, tracing of requests, magic anomaly detection, accounting, SLA reporting

simple single node w/ local storage for a few weeks

over

horizontal scaling, clustering, multitenancy

configuration files

over

Web UI, user management

pulling data from single processes

over

pushing data from processes, aggregation on nodes

NoSQL query & data massaging multidimensional data everything as float64

over

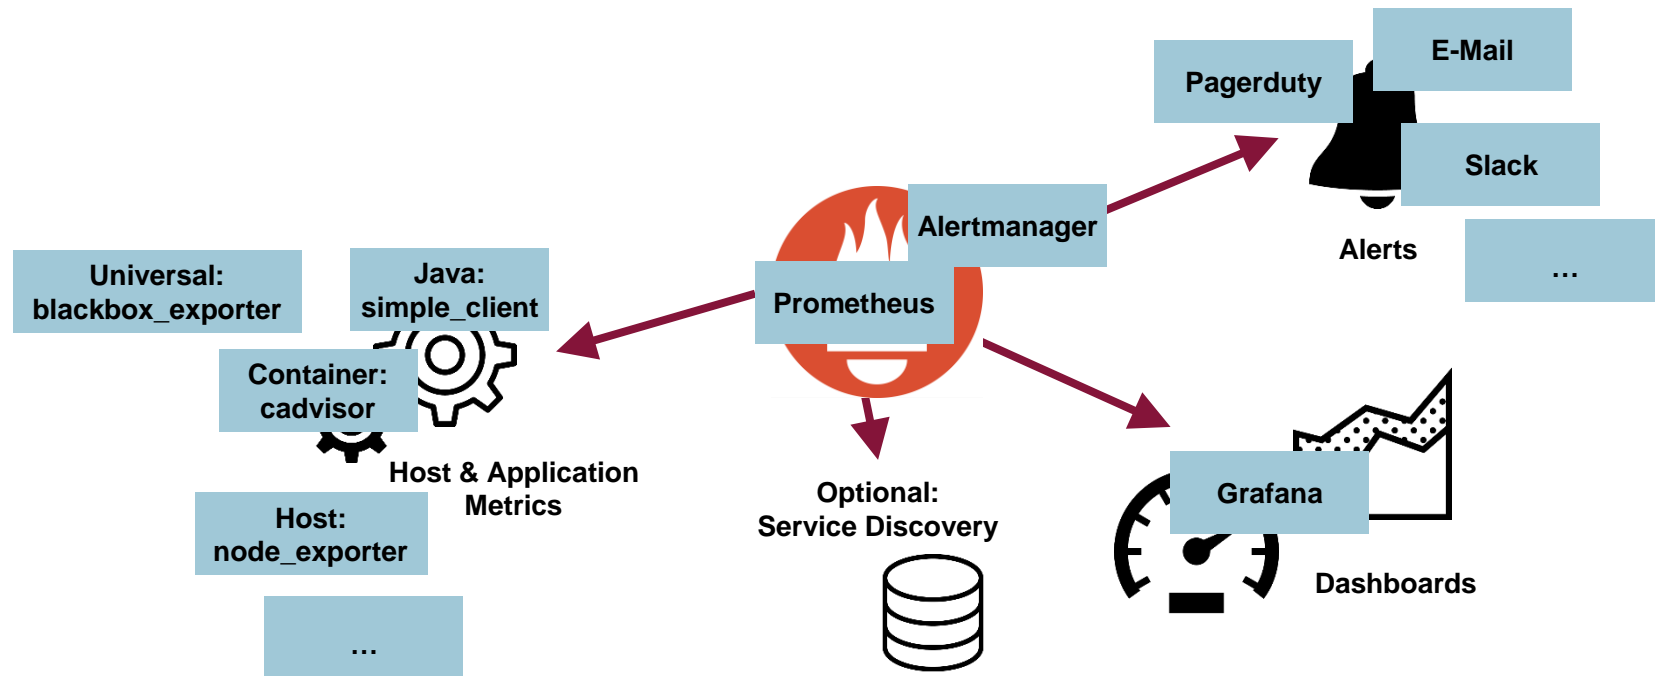
point-and-click configurations, data silos, complex data types

1. PromCon 2016: Prometheus Design and Philosophy - Why It Is the Way It Is - Julius Volz <https://youtu.be/4DzoajMs4DM>
/ <https://goo.gl/1oNaZV>

Monitoring for Developers with Prometheus and Grafana

- 1 Prometheus Manifesto
- 2 **Setup**
- 3 How to...
- 4 Prometheus works for Developers (and Ops)

Setup



Technical Building Blocks






```
scrape_configs:  
  - job_name: 'node-exporter'  
    scrape_interval: 5s  
    static_configs:  
      - targets: ['172.17.0.1:9100']
```

Setup

CPU Metric as exported by the Node Exporter

```
# HELP node_cpu Seconds the cpus spent in each mode.
# TYPE node_cpu counter
node_cpu{cpu="cpu0",mode="guest"} 0
node_cpu{cpu="cpu0",mode="idle"} 4533.86
node_cpu{cpu="cpu0",mode="iowait"} 7.36
...
node_cpu{cpu="cpu0",mode="user"} 445.51
node_cpu{cpu="cpu1",mode="guest"} 0
node_cpu{cpu="cpu1",mode="idle"} 4734.47
...
node_cpu{cpu="cpu1",mode="iowait"} 7.41
node_cpu{cpu="cpu1",mode="user"} 576.91
...
```

Multidimensional Metric as stored by Prometheus



Calculations based on metrics

Metric:

node_cpu: Seconds the CPUs spent in each mode (Type: Counter).

What percentage of a CPU is used per core?

```
1 - rate(node_cpu{mode='idle'} [5m])
```



What percentage of a CPU is used per instance?

```
avg by (instance) (1 - rate(node_cpu{mode='idle'} [5m]))
```

Monitoring for Developers with Prometheus and Grafana

- 1 Prometheus Manifesto
- 2 Setup
- 3 How to...**
- 4 Prometheus works for Developers (and Ops)

Information about your containers

Presented by: cadvisor

RAM Usage per container:

Variable: `container_memory_usage_bytes`

Expression: `container_memory_usage_bytes{name=~'.+',id=~'/docker/.*'}`

CPU Usage per container:

Variable: `container_cpu_usage_seconds_total`

Expression: `rate(container_cpu_usage_seconds_total [30s])`
`irate(container_cpu_usage_seconds_total [30s])`
`sum by (instance, name) (irate(container_cpu_usage_seconds_total{name=~'.+'} [15s]))`

Information about your JVM

Presented by: Java simple_client

RAM Usage of Java VM:

Variable: `jvm_memory_bytes_used`

Expressions: `irate(container_cpu_usage_seconds_total [30s]) sum
by (instance, job) (jvm_memory_bytes_used) sum by
(instance, job) (jvm_memory_bytes_committed)`

CPU seconds used by Garbage Collection:

Variable: `jvm_gc_collection_seconds_sum`

Expression: `sum by (job, instance) (irate(jvm_gc_collection_seconds_sum [10s]))`

Information about your JVM

Add a Configuration to Spring Boot to serve standard JVM metrics using /prometheus actuator endpoint.

```
@Configuration
@EnablePrometheusEndpoint
public class ApplicationConfig {

    @PostConstruct
    public void metrics() {
        DefaultExports.initialize();
        /* ... */
    }
}
```

Information about your Application Metrics

Presented by: Java simple_client and Spring

Timings of a method call:

Java Annotation: @PrometheusTimeMethod(name = "example", help = "...") Variables:

example_count

example_sum

Information about your Application Metrics

Add a Configuration to collect Prometheus timings from Annotations.

```
@Configuration
@EnablePrometheusTiming
public class MetricsApplicationConfig {

    /* ... */
}
```


Information

about your Application Metrics

Add @PrometheusTimeMethod annotations to any method of any Bean to collect metrics

```
@Component
public class RestEndpoint {

    @Path("countedCall")
    @GET
    @PrometheusTimeMethod(name = "example", help = "...")
    public Response countedCall() throws InterruptedException {
        /* ... */
        return Response.ok("ok").build();
    }

}
```

Information about your External Interfaces

Presented by: Java simple_client, Hystrix/Spring

Hystrix Metrics:

Java Annotation: @HystrixCommand

Variables: hystrix_command_total {command_name="externalCall", ...}
 hystrix_command_error_total {command_name="externalCall", ...}

Expressions: histogram_quantile(0.99,
 rate(hystrix_command_latency_execute_seconds_bucket[1m]))

External Interfaces – Hystrix Metrics

Register the Hystrix Publisher and add @HystrixCommand for resilience and timing of external calls.

Information about your

```
HystrixPrometheusMetricsPublisher.register();
```

```
@Component
public class ExternalInterfaceAdapter {

    @HystrixCommand(commandKey = "externalCall", groupKey = "interfaceOne")
    public String call() {
        /* ... */
    }
}
```

Information about your Spring Servlet Container

Presented by: your own Java metric provider

Tomcat Connector:

Java Class: Write your own: TomcatStatisticsCollector

Variables: tomcat_thread_pool_current_thread_count
tomcat_thread_pool_current_threads_busy

Tomcat DB Connection Pool:

Java Class: Write your own: DatasourceStatisticsCollector

Variables: tomcat_datasource_active tomcat_datasource_idle
tomcat_datasource_max_idle

How to...

Information about your

Spring Servlet Container

```
public class DatasourceStatisticsCollector extends Collector {  
  
    /* ... */  
  
    @Override  
    public List<MetricFamilySamples> collect() {  
        /* ... */  
        result.add(buildGauge("active", "number of connections in use",  
                                labelNames, labelValues, tomcatDS.getActive()));  
        return result;  
    }  
  
}
```

```
new DatasourceStatisticsCollector(dataSource).register();
```

Information about your

Information about your Vert.x application

Presented by: Java Simple Client for Vert.x

Internal Event Bus:

Variables: vertx_eventbus_messages_sent_total
 vertx_eventbus_messages_pending
 vertx_eventbus_messages_delivered_total
 vertx_eventbus_messages_reply_failures_total

HTTP Server metrics:

Variables: vertx_http_servers_..._requests_count
 vertx_http_servers_..._open_netsockets

Vert.x application

```
// During Setup
vertx = Vertx.vertx(new VertxOptions().setMetricsOptions (
    new DropwizardMetricsOptions ()
        .setRegistryName ("vertx")
        .addMonitoredHttpClientEndpoint (
            new Match().setValue (".*") .setType (MatchType.REGEX) )
        .setEnabled (true)
));

DefaultExports.initialize ();
new DropwizardExports (SharedMetricRegistries.getOrCreate ("vertx")) .register ();

// When starting up Routes and a HTTP Server
final Router router = Router.router (vertx);
router.route ("/metrics") .handler (new MetricsHandler ());
```

Information about your

Federation of Prometheus

Any Metric can be exported to other Prometheus instances

`http://localhost/prometheus/federate?match[]={job=%22prometheus%22}`

Alerting with Prometheus

Any expression can be used for alerting

alert: HDD_Alert_warning

```
(1 - node_filesystem_free{mountpoint=~".*"} / node_filesystem_size{mountpoint=~".*"}) * 100 > 70
```

expr: for: 5m labels:

severity: warning

annotations: summary: High disk usage on {{ \$labels.instance }}: filesystem {{ \$labels.mountpoint }} more than 70 % full.

Setup of the Environment

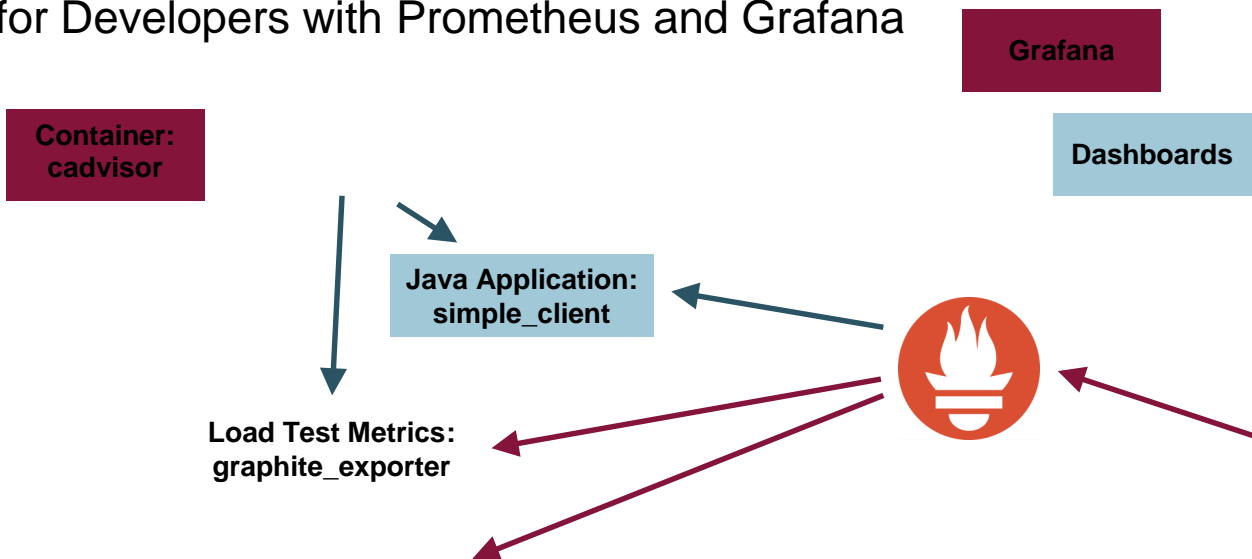
Technical Building Blocks for Load Testing

Load Test:

Purple (including Prometheus): Provided as infrastructure in a testing environment

Blue: Setup and maintained by product team (developers/testers)

Monitoring for Developers with Prometheus and Grafana



-
- 1 Prometheus Manifesto
 - 2 Setup
 - 3 Howto...
 - 4 Prometheus works for Developers (and Ops)**

What to expect

Lessons learned

The approach worked well for us to pass the load tests:

- Load Tool metrics correlated with application and infrastructure metrics
- Inter-application communication captured by Hystrix
- Self-service functionality for product teams to add applications and metrics

... but to use the instrumentation also in production create awareness:

- Exported metrics should follow Prometheus naming conventions
- Collector for Dropwizard Metrics can't fill HELP text of metrics
- Counters and averages vs. histograms, summaries and percentiles
- Consistent use of USE Method (utilization – saturation – errors) or RED Method (rate – errors – duration) for metrics

1. <http://www.brendangregg.com/usemethod.html>

2. <https://www.weave.works/blog/prometheus-and-kubernetes-monitoring-your-applications/>

Prometheus works for Developers (and Ops)

Prometheus is “friendly tech” in your environment

Team friendly

- Every team can run its own Prometheus instance to monitor their own and neighboring systems
- Flexible to collect and aggregate the information that is needed

Coder and Continuous Delivery friendly

- All configurations (except dashboard) are kept as code and are guarded by version control
- Changes can be tested locally and easily staged to the next environment

Simple Setup

- Go binaries for *prometheus* and *alertmanager* available for major operating systems
- Client libraries for several languages available (also adapters to existing metrics libraries)
- Several existing exporters for various needs

Links

Prometheus:

<https://prometheus.io>

Java Simple Client

https://github.com/prometheus/client_java

Hystrix

<https://github.com/Netflix/Hystrix>

Prometheus Hystrix Metrics Publisher

<https://github.com/ahus1/prometheus-hystrix>

Dropwizard Metrics

<http://metrics.dropwizard.io>

Prometheus Design and Philosophy

- Why It Is the Way It Is

<https://youtu.be/4DzoajMs4DM>

<https://goo.gl/1oNaZV>

CAdvisor

<https://github.com/google/cadvisor>

